

Applying Principal Component Analysis for Noise Reduction in Image Classification

Andrew Duenner
duenner@mit.edu

Alex Kimn
akimn@mit.edu

May 19, 2018

Contents

1 Introduction

As computer vision becomes more ubiquitous, expectations for speed and robustness continue to grow. This report focuses on the application of Principal Component Analysis (PCA) as a pre-processor to improve the robustness of image classifiers with respect to noisy data. Specifically, we consider Convolutional Neural Networks (CNNs) trained on the MNIST dataset of handwritten characters. Methods for calculating principal components using singular value decomposition (SVD), eigenvector decomposition, and partial least squares minimization are introduced.

Each method is implemented in Python and compared on the basis of *Normalized Run-time*, *Loss*, *Error Rate*, and *Mean Squared Error (MSE)*. Methods for calculating principal components are also compared on the basis of asymptotic operation count with respect to data set size, image size, and number of principal components used. Finally, the advantages and disadvantages of each method applied to image noise reduction are discussed and conclusions are presented.

2 Overview of Principal Components

Principal component analysis is a mathematical transformation $X = TP^T$ that transforms a data set $X \in \mathbb{R}^{n \times p}$ consisting of n samples and p variables with unknown correlation to a set of uncorrelated orthogonal loading vectors P that map X to a matrix T of principal component scores. The transformation operates such that the first vector of P corresponds to the direction of the original data set with the largest variance [jolliffe1986principal]. This property of principal components is useful for noise reduction in a data set as it allows for the correlated data to be separated from the uncorrelated data. X can be approximated as \hat{X} by using the first k principle components such that $\hat{X} = T_k P_k^T$ with $X = \hat{X} + E$ and E an error term representing the remainder of X in the directions with the lowest correlation.

In **Figure ??** on the following page, we present an example of PCA applied to the MNIST dataset. Each of the individual images is the result of projecting one noised digit onto the first n principal components, which are extracted across the entirety of the noised dataset.

Effect of Varying $n_{components}$ on PCA Output




Figure 1: Result of projecting a noisy image representing a 6 onto various numbers n of principal components

3 Algorithms

3.1 NIPALS

The nonlinear iterative partial least squares (NIPALS) algorithm is an iterative algorithm that has the advantage of calculating principal components consecutively and with arbitrary accuracy. These attributes are desirable when the number of required principal components is uncertain or variable and the requirements on accuracy are somewhat relaxed. As such, it is an algorithm well-suited for pre-processing noised images quickly.

The NIPALS algorithm consists of an outer loop that calculates the i^{th} next consecutive principal component. The outer loop begins with a guess of the i^{th} score vector. The initial value for t_i may be set to the i^{th} column of the data set X_i where i is the column index of the principal component being calculated [risvik2007principal]. Next the inner product of t_i is taken with respect to itself as λ_i which represents the square of the magnitude of t_i . An inner loop next iteratively calculates the i^{th} loading vector p_i , score vector t_i and eigenvector λ_i . Within the inner loop p_i is calculated by projecting the transpose of the data set (X^T) onto the scores t_i . Vector p_i is then normalized to prevent overflow. The scores vector t_i is updated to be the projection of X onto p . A new λ_{new} is calculated as the inner product of t with itself. If $|\lambda - \lambda_{new}| < tol$ the iterative inner loop breaks whereas otherwise λ is set to λ_{new} . After a number of iterations t_i converges to the component score vector, p_i converges to the component loading vector, and λ_i converges to the i^{th} largest eigenvalue of X . After the inner loop converges, the outer product of t_i and p_i is subtracted from X and the outer loop is repeated. The code snippet in Listing ?? below demonstrates an implementation of the NIPALS method in Python.

The computational complexity of the NIPALS algorithm is dominated by the matrix vector product operations used to project X onto t . Projecting X^T onto t and X onto p are both $\mathcal{O}(\mathbf{mn})$ for a data set with \mathbf{m} samples and \mathbf{n} variables. Both of these operations take place within the inner loop of k iterations to convergence and an outer loop of \mathbf{c} principal components. The entire algorithm is then $\mathcal{O}(k\mathbf{c} \times \mathbf{mn})$ where \mathbf{m} is images, \mathbf{n} is pixels per image, \mathbf{c} is number of principal components and k is number of inner loop iterations necessary to reach desired accuracy.

Code Listing 1: NIPALS Algorithm

```
# Nonlinear Iterative Partial Least Squares
@jit(nopython=True) # Numba
def NIPALS(X, scores, loadings, lambdas, n_components,
           threshold=1e-6, max_iter=200):

    # Iterate over each calculated component
    for i in range(n_components):

        t = X[:, i]
        lambda_ = np.dot(t, t)

        # Iterate until convergence
        for j in range(max_iter):

            # Compute loadings
            p = np.dot(X.T, t)
            p /= lambda_
            loadings[:, i] = p

            p /= np.linalg.norm(p)

            # Project X onto p to find score vector t
            t = np.dot(X, p)
            t /= np.dot(p, p)
            scores[:, i] = t

            lambda_n = np.dot(t, t)

            # Add score vector to matrix of score vectors
            lambdas[i] = lambda_n

            # Check for convergence
            diff = np.abs(lambda_ - lambda_n)

            if diff < threshold:
                break

            lambda_ = lambda_n

    # Update X
    X -= np.outer(t, p)
```

3.2 Full SVD

Singular Value Decomposition is a commonly implemented method for calculating principal components [oliphant2009numpy] [scikit-learn]. For a mean-centered data set $X \in \mathbb{C}^{m,n}$, the SVD transformation satisfies the equality $X = U\Sigma V^*$ where $U \in \mathbb{C}^{m,n}$ is a unitary matrix representing the left singular vectors of X , $V \in \mathbb{C}^{n,n}$ is a unitary matrix representing the right singular vectors of X , and $\Sigma \in \mathbb{C}^{m,n}$ is a diagonal matrix of the singular values of X [trefethen1997numerical]. The principal component score vectors t_i can be calculated as $U\sigma$ and the loading vectors are VX [madsen2004singular].

SVD decomposition was chosen as a benchmark algorithm to understand the implications of calculating only a few principal components. A common algorithm for calculating singular value decomposition of a matrix is the gebrd algorithm in LAPACK. This algorithm is of order $\mathcal{O}((4/3) * n^2 * (3m - n))$ [intelmath] and calculates all of the singular values and singular vectors simultaneously. Thus the NIPALS algorithm should be superior to SVD in the event that $(4/3) * n^2 * (3m - n) > (kc \times mn)$.

An implementation of principal component analysis by SVD using Python and the numpy library is given below in code listing ??.

Code Listing 2: Calculating Principal Components with SVD

```
@jit(nopython=True) # Numba
def SVD(X, scores, loadings, lambdas, n_components):

    u, s, vt = np.linalg.svd(X, full_matrices=False)
    m = X.shape[0]

    loadings[:, :] = (vt.T)[:, 0:n_components]
    scores[:, :] = np.dot(u, np.diag(s))[:, 0:n_components]
    lambdas[:, :] = np.dot(s, s)
    lambdas /= (m - 1)
```

3.3 Eigenvector Decomposition (Simultaneous Iteration)

Eigenvector decomposition provides yet another method for calculating principal components. An iterative method for simultaneously calculating the the first few eigenvalues and eigenvectors was selected as a benchmark algorithm to provide a comparative basis for iteratively calculating principal components sequentially as in NIPALS versus calculating them simultaneously.

Calculating principle components with eigenvector decomposition starts with calculation of the covariance C of the data set represented such that the rows are pixels and columns are samples.

$$C = \frac{1}{n-1} X^T X$$

Calculating the covariance matrix is of $\mathcal{O}(\mathbf{mn}^2)$ and is a significant component of the algorithms overall computational efficiency.

Eigenvector decomposition is a transformation of the form $C = Q\lambda Q'$ where Q are the eigenvectors of C and λ are the eigenvalues. Simultaneous iteration solves for the k largest eigenvalue eigenvector pairs simultaneously by repeatedly multiplying the covariance matrix $C \in \mathbb{R}^{m,n}$ multiplied by a matrix $Q^{n,c}$ which is initialized with random values. Matrix Q converges to the c largest eigenvectors of the covariance matrix C [trefethen1997numerical]. In order to prevent the largest eigenvector from dominating all columns in Q , Matrix A is reorthogonalized every iteration using the modified gram-schmidt procedure adapted from [srinivasan]. Q_0 is then set to Q and the value of Q is tested for convergence. Convergence conditions are considered satisfied when the orthogonal projection of Q onto Q_0 is sufficiently small as discussed in [arbenz2012lecture].

$$\|(I - Q_k Q_k^T) Q_{k-1}\| < tol$$

The principal component score matrix T and loading vector P can be constructed from the covariance C , Eigenvectors Q and Eigenvalues R as follows.

$$T = QC$$

$$P = V$$

The outer loop of the simultaneous power iteration algorithm is dominated in computational complexity by the multiplication of C with Q which is $\mathcal{O}(\mathbf{mnc})$ and the inner loop in which modified gram schmidt is used to orthogonalize A is of $\mathcal{O}(2\mathbf{mc}^2)$. The total complexity is then approximately $\mathcal{O}((2\mathbf{mc}^2 + \mathbf{mnc})k)$ and is highly sensitive to the number of eigenvectors chosen for simultaneous computation. Code listing ?? below shows a numpy implementation of simultaneous iteration with orthogonalization at each step by MGS.

Code Listing 3: Calculating Principal Components with Eigenvalue Decomposition via Simultaneous Power Iteration


```

@jit(nopython=True) # Numba
def SI(X, Q, Q_o, R, I, threshold=1e-6, max_iter=200):

    for k in range(max_iter):

        A = operator.matmul(C, Q)

        Q_o[:] = Q

        _MGS(A, Q, R)

        delta = operator.matmul(I - np.dot(Q, Q.T), Q_o)

        if np.linalg.norm(delta) <= threshold:
            break

@jit(nopython=True) # Numba
def MGS(A, Q, R):

    (m, n) = A.shape

    for k in range(n):

        R[k, k] = np.sqrt(np.dot(A[:, k], A[:, k]))
        Q[:, k] = A[:, k] / R[k, k]

        for j in range(k + 1, n):

            R[k, j] = np.dot(Q[:, k].T, A[:, j])
            A[:, j] -= R[k, j] * Q[:, k]

```

3.4 GPU Paralellization

Consumer grade graphics processing units have been released with the capability to perform parallel computation of floating point arithmetic at high speed. In this study GPU hardware was used for the neural network classifier and so it was a natural extension to implement the image preconditioner in GPU-executable code as well. It was hypothesized that at least some reduction in computation time could be achieved by taking advantage of parallelization of matrix operations.

PyTorch was used as a high-end language to execute the NIPALS algorithm on a consumer-grade gaming GPU. Conversion from numpy to PyTorch proved to be a trivial development exercise and

the resulting algorithm ran significantly faster on the GPU when compared to execution on a CPU. Figure ?? shows a comparison of the CPU and GPU implementations for a range of image sizes with. The number of components calculated, $n_{components}$, is set to 5 and full images were used ($n_{pixels} = 784$).




Figure 2: A flow chart of the bench-marking process

The data shows that runtime on the GPU is a factor of 4-5 times less than that of the CPU. Note that the GPU advantage is dependent on the number of images and for smaller datasets the CPU implementation is faster. This is likely due to the overhead associated with memory bandwidth and transfer between RAM on the CPU and GPU.

4 Methodology

In order to assess the relative performance of each of the algorithms, we measure four separate metrics: *Normalized Run-time*, *Loss*, *Error Rate*, and *Mean Squared Error (MSE)*. **Figure ??** below provides an overview of how this is achieved.

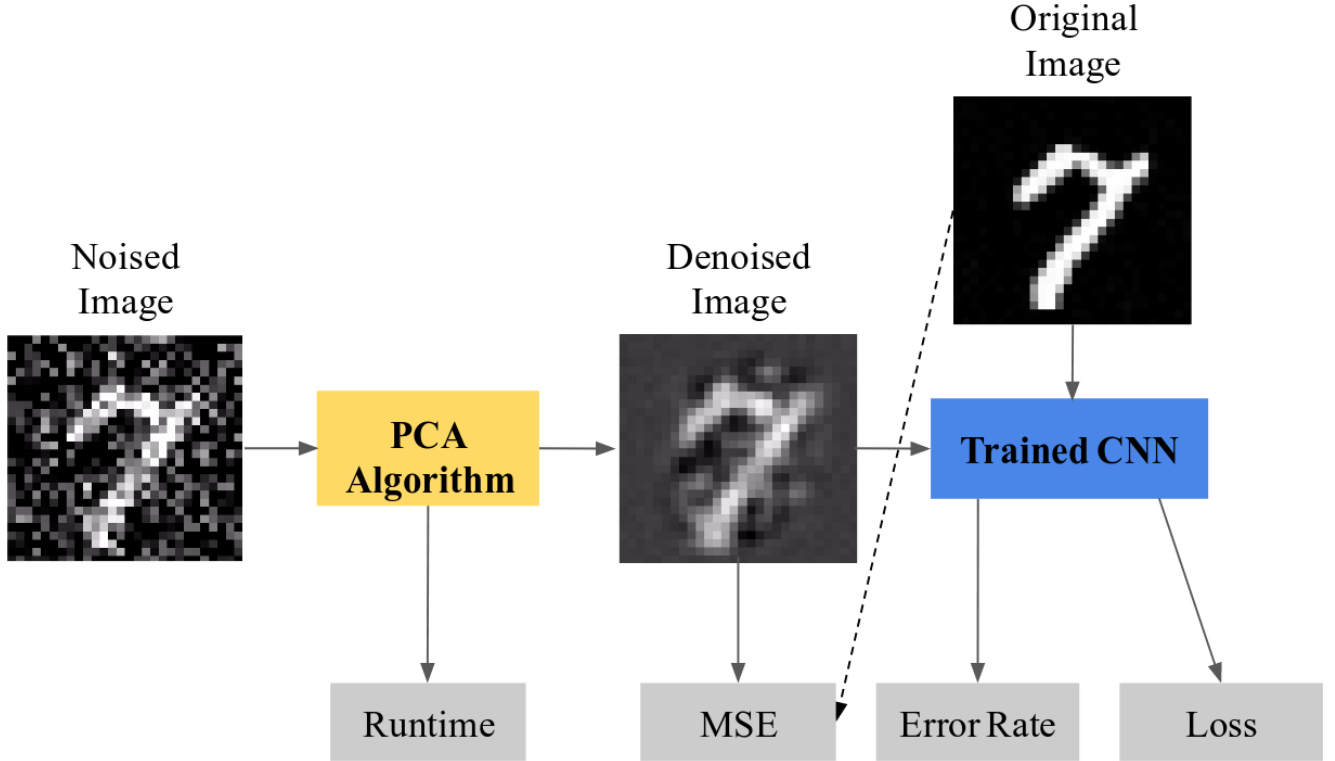


Figure 3: A flow chart of the bench-marking process

A fair assessment of the relative cost of running each algorithm is rather difficult to obtain. While we implement all three using only the Python library *NumPy*, its built-in functions use a variety of different back-ends. Therefore, we wrap each of the *NumPy* implementations inside *Numba*, so that all three algorithms are compiled by LLVM and run as machine code.

Moreover, since neither *NumPy* nor *Numba* offer methods to directly determine the number of used floating point operations, we resort to using *Normalized Run-time* instead as a proxy for cost.

Of course, we cannot guarantee that raw runtimes across methods are proportional to relative cost. However, given the unified *Numba* backend, we should expect that the runtimes of different instances of the same method are proportional to computational complexity. Hence, if these runtimes are then normalized at some fixed value of the independent variable, we can compare their relative trajectory across that variable. Using this, we evaluate the computational complexities of the three algorithms on two values of dataset size, the number of images, n_{images} , and the number of pixels per image, n_{pixels} , as well as the number of components extracted, $n_{components}$.

In addition to the algorithms' running costs, we also measure the quality of their output. In the context of the original problem, higher quality outputs should result in an improvement in the classifier's discriminatory ability. Since this quality depends on $n_{components}$, we test each PCA algorithm on a set of 20 $n_{components}$ values ranging from 10 to 680. Additionally, we apply each of these 20 variants to 8 different noised datasets (with mean noise mask magnitude ranging from $\mu_{noise} = 10$ to 80), resulting in a total of 160 different de-noised outputs for each algorithm. From there, we run these outputs through a set of five Convolutional Neural Network (CNN) classifiers, each trained using *PyTorch* on the same un-noised training data, and record the average *Loss* and average *Error Rate* that is observed.

Finally, noting that the trained classifiers are all unique and thus introduce additional variance into the problem, we select an objective evaluation of image quality. In particular, for each noised image N_i in a noised dataset (where i is its index in the set), we set the image it was derived from, O_i as the ground truth and determine the mean squared error of the de-noised image, D_i with respect to it. These values are then averaged across the entire set to obtain the final value of *MSE*, as shown below:

$$MSE = \frac{1}{n_{images}} \frac{1}{n_{pixels}} \sum_{i=1}^{n_{images}} \sum_{j=1}^{n_{pixels}} (D_i[j] - O_i[j])^2$$

As with *Loss* and *Error Rate* we calculate *MSE* for each of the 8 magnitude of noise and several different values for $n_{components}$.

5 Algorithm Benchmark Results

5.1 Normalized Runtime

Before analyzing the results obtained from testing the PCA algorithms, we analyze the general asymptotic time complexity of each of the algorithms as they are implemented in *Numba*. Please refer back to sections 3.1, 3.2, and 3.3 for the raw code used for each algorithm.

5.1.1 Asymptotic Complexity

We note that the computational complexity of the three PCA algorithms is primarily determined by three values: $\mathbf{m} == n_{images}$, $\mathbf{n} == n_{pixels}$, and $\mathbf{c} == n_{components}$. Given these variables, we can estimate the asymptotic complexity of each as follows:

NIPALS: The operation count of the inner loop is dominated by the calculation of the two dot products involving the data array X . Knowing that X has size (\mathbf{m}, \mathbf{n}) , we can deduce that both of these function calls are $\mathcal{O}(\mathbf{mn})$. Then, taking into consideration the two loops, we find that the NIPALS method is $\mathcal{O}(k \times \mathbf{c} \times \mathbf{mn})$, where $k \leq max_iter$ represents the number of inner loop iterations needed for convergence.

Full SVD: The method's cost is dominated by producing the SVD of the original data, which is size (\mathbf{m}, \mathbf{n}) . Hence, we expect this method to have time complexity $\mathcal{O}(\mathbf{mn} \times \min[\mathbf{m}, \mathbf{n}])$. Note that since the algorithm computes all the principal components by default, the runtime has no dependence on \mathbf{c} .

Simultaneous Iteration: The majority of the computation for each iteration is done to complete a Modified Gram-Schmidt (MGS) orthogonalization. Given that the inputted matrix X has dimensions (\mathbf{n}, \mathbf{n}) , we determine that the complexity of each inner loop within the MGS iterations is $\mathcal{O}(\mathbf{n})$. Since both the inner and outer loops of MGS are called \mathbf{c} times, each MGS call has time complexity $\mathcal{O}(\mathbf{nc}^2)$. Note that since the size of the input to the algorithm has no relation to the number of images inputted, the runtime has no dependence on \mathbf{m} . Finally, as with NIPALS above, Simultaneous Iteration also loops at $k \leq max_iter$ times until convergence. Hence we find that this algorithm is $\mathcal{O}(k \times \mathbf{nc}^2)$.

To confirm this analysis, we vary all three of \mathbf{m} , \mathbf{n} , and \mathbf{c} and measure the resulting *Normalized Runtime* of each of the algorithms as they are applied to a set of 70000 images noised with a mask

of magnitude 30.

5.1.2 Runtime and $\mathbf{m} == n_{images}$

The first variable we examine is $\mathbf{m} == n_{images}$. To do so, we use full images (i.e. $n_{pixels} = 784$), and fix $n_{components}$ to be 50. We then vary n_{images} from 100 to 1000 and run 15 trials for each of the algorithms, using the median as the representative value. We randomly select images from the dataset to ensure that representative values of k are obtained for both **NIPALS** and **Simultaneous Iteration**. Finally, in

The results of this test are shown in a raw and normalized form below in **Figures ??** and **??**, respectively. Note that in the normalized graph, all three methods are chosen to have runtime 1 at $n_{images} = 100$.




Figure 4: Unnormalized runtime behavior of PCA algorithms when n_{images} is varied.

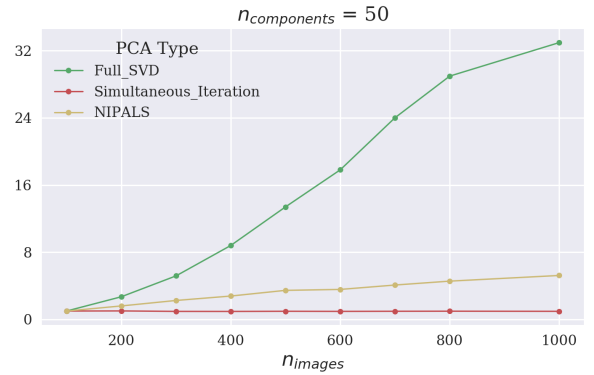


Figure 5: Normalized runtime behavior of PCA algorithms when n_{images} is varied.

Given the time complexities of the algorithms, we expect to see a linearly increasing runtime for both **NIPALS** and **Full SVD**. And indeed, this trend is clearly visible on the normalized graph. Moreover, our expectation of constant runtime for the **Simultaneous Iteration** algorithm is also confirmed by the data.

5.1.3 Runtime and $n == n_{pixels}$

Next, we show how varying the value of $n == n_{pixels}$ affects the time cost of the algorithms. Specifically, we fix n_{images} to 5000 and $n_{components}$ to 5 and vary n_{pixels} from 20 to 320. We run 25 trials for each of the algorithms, selecting the median value as the representative value. As before, we randomly select the images (and also pixels) used in each individual trial to ensure that we obtain representative values of k .

The results of this test are shown in a raw and normalized form below in **Figures ??** and **??**, respectively. Note that in the normalized graph, all three methods are chosen to have runtime 1 at $n_{pixels} = 20$.

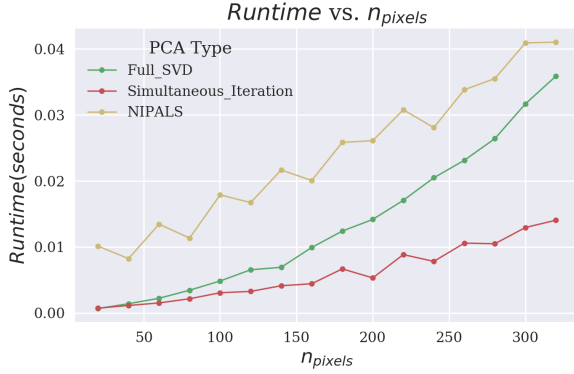


Figure 6: Unnormalized runtime behavior of PCA algorithms when n_{pixels} is varied.




Figure 7: Normalized runtime behavior of PCA algorithms when n_{pixels} is varied.

This specific setup ensures that $n_{images} > n_{pixels}$. Hence from the complexity analysis above, we predict **Full SVD** to have quadratic complexity in n_{pixels} , which the data appears to confirm. Moreover, our expectations of linear runtimes for both **NIPALS** and **Simultaneous Iteration** are also supported by the graphs.

5.1.4 Runtime and $c == n_{components}$

Finally, we analyze $n_{components}$. Here, we fix n_{images} to be 1000 and again use full images (i.e. $n_{pixels}=784$). We run 15 trials, randomly selecting the images and taking the median runtime as the representative.

The results of this test are shown in a raw and normalized form below in **Figures ??** and **??**, respectively. Note that in the normalized graph, all three methods are chosen to have runtime 1 at $n_{components} = 10$.




Figure 8: Unnormalized runtime behavior of PCA algorithms when n_{images} is varied.




Figure 9: Normalized runtime behavior of PCA algorithms when n_{images} is varied.

The data confirms the predicted linear relation between the cost and $n_{components}$ for **NIPALS**. The quadratic complexity of **Simultaneous Iteration** is also clearly visible in the normalized graph. Finally, the data confirms that the cost of **Full SVD** does not depend at all on the number components extracted.

5.1.5 Overview

In general, the above assessments confirm the theoretical complexities of three three PCA algorithms. However, it follows that the relationship between computational and the data parameters differs for all three algorithms. And if we extrapolate the complexity relationships observed to additional datasets and problems, we should find that the computationally cheapest PCA algorithm may vary.

In particular, we would expect **NIPALS** to excel in problems where $n_{components}$ is low and the total data size of the problem is sufficiently large. However, if the number of unique data samples is sufficiently large, we would expect **Simultaneous Iteration** to be the most efficient. Finally, for problems like image de-noising, where the the value of $n_{components}$ is high and the dataset size relatively small, we should expect to see that **Full SVD** be the most efficient.

Indeed, as observed in the following sections, de-noising the MNIST dataset often required up to 100 principal components to achieve the best noise reduction, and as shown by the unnormalized graphs above, the raw runtime for **Full SVD** was often only a fraction of that of the other two methods.

5.2 Mean Squared Error

The obtained values of MSE for both the lowest noise ($\mu_{noise} = 10$) and highest noise ($\mu_{noise} = 80$) datasets are shown below in **Figures ??** and **??**, respectively.




Figure 10: The relationship between MSE and $n_{components}$ for the dataset with $\mu_{noise} = 10$




Figure 11: The relationship between MSE and $n_{components}$ for the dataset with $\mu_{noise} = 80$

The three different PCA algorithms appear to perform nearly identically on the least noisy dataset, but for the noisiest dataset, there is a non-trivial difference between the algorithms. In particular, it appears that the de-noised images from the **Full SVD** algorithm are consistently more distant from the ground truth than the de-noised images from the other two variants.

Additionally, for the least noisy dataset, as $n_{components}$ increases, the value of MSE generally decreases. However, for the dataset with $\mu_{noise} = 80$ and the other high noise datasets, there appears to be an optimal value for $n_{components}$. That is, there is a definite minimum value of $n_{components}$ that minimizes the value of MSE . Therefore, we can conclude that the choice of $n_{components}$ not only has a significant impact on the algorithms' run-times but also has a large influence on the quality of images they produce.

5.3 Loss

As in the previous section, the obtained values of *Loss* for both the lowest noise ($\mu_{noise} = 10$) and highest noise ($\mu_{noise} = 80$) datasets are shown below in **Figures ??** and **??**, respectively.




Figure 12: The relationship between *Loss* and $n_{components}$ for the dataset with $\mu_{noise} = 10$




Figure 13: The relationship between *Loss* and $n_{components}$ for the dataset with $\mu_{noise} = 80$

Like before, we see that the different PCA variants perform nearly identically for the lowest noise dataset. However, for the highest noise dataset something is quite different. Specifically, whereas before the **Full SVD** algorithm performed worse than its peers on the *MSE* metric, on the *Loss* metric, it appears to consistently perform the best out of the three. The reason for both of these discrepancies is unclear, though it may result from the small non-zero value added to the mean of both the **NIPALS** and **Simultaneous Iteration** to prevent division by zero-norm vectors. In any case, we believe that the relative differences are not significant enough to have a noticeable impact on the algorithms' performance.

Additionally, a curious difference in the optimal value of $n_{components}$ between the *Loss* and *MSE* metrics appears. In the case of *Loss*, it appears that the optimal value is somewhat smaller and the local minimum somewhat steeper. This may result from the fact that classifiers are quite sensitive to the noise that a larger number of $n_{components}$ retains.

5.4 Error Rate

The data concerning the *Error Rate*, again for the least and most noisy datasets are shown below in **Figures ??** and **??**, respectively.




Figure 14: The relationship between *Error Rate* and $n_{components}$ for the dataset with $\mu_{noise} = 10$




Figure 15: The relationship between *Error Rate* and $n_{components}$ for the dataset with $\mu_{noise} = 80$

These results roughly mirror those concerning *Loss*. This is hardly surprising, as the two metrics are inextricably linked by the classifiers. The main point of interest here is the fact that the relative discrepancy that was seen in the *Loss* graphs is somewhat smaller in the *Error Rate* graphs. This probably results from the fact that small changes in loss do not necessarily result in the re-classification of a sample.

6 Application to Classifier

In the **analysis** of the *MSE* data, we noted that the ideal value of $n_{components}$ appears to vary with the magnitude of noise applied. We also noted that this value appears to change depending on the metric used.

We show this below for all 8 noise levels, using both *MSE* and *Loss*, in **Figures ??** and **??**, respectively. Note that in these figures, the metrics are averaged over both the 5 classifiers and the 3 different PCA methods.




Figure 16: The relationship between *MSE* and $n_{components}$ for all 8 noised datasets

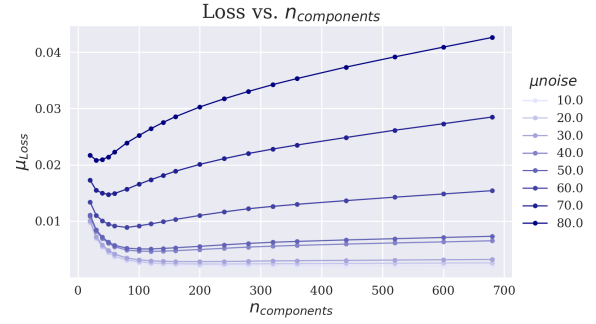


Figure 17: The relationship between *Loss* and $n_{components}$ for all 8 noised datasets

From these data, we can again determine that the optimal value of $n_{components}$ decreases as noise increases. And as before, we note that the curves for *Loss* (and also for *Error Rate*) reach their minimum earlier and are steeper than those relating to *MSE*.

Nevertheless, since the end goal of using these PCA algorithms is to improve the performance of the classifier, we utilize the optimal value of $n_{components}$ with respect to *Loss* in order to select the final PCA algorithm. In particular, for each noise value and each classifier, we select the value of $n_{components}$ that resulted in the lowest *Loss* value.

We give an example of how the selected PCA algorithms impact the classifiers' performances below in **Figures ??** and **??** below, which show the average *Loss* and *Error Rate* respectively of the classifiers when each of the selected PCA algorithms is applied. The graphs also show the average performance of the classifiers on the raw noised data for comparison.




Figure 18: The average *Loss* of the classifiers when a PCA pre-processor is applied to the noised data




Figure 19: The average *Error Rate* of the classifiers when a PCA pre-processor is applied to the noised data

The use of any of the three PCA algorithms results in roughly the same result in the classifier. Specifically, for lower values of noise, it appears that the use of a PCA pre-processor does not help the classifier at all; in fact, both the average *Loss* and *Error Rate* of the classifiers increases when the noise is low. However, at sufficiently high values of noise, it appears that using a PCA pre-processor has a significant positive impact on classifier performance. Indeed for the most noised data available, the use of any of the PCA algorithms resulted in a $> 20\%$ increase in the accuracy of the classifier. And we surmise that if still more noise were to be added to the data, an even more significant difference could be observed. Hence, it appears that PCA can function as an image denoiser, particularly for data with extremely high levels of noise.

7 Conclusion

Our results demonstrate that Principal Component Analysis (PCA) can be an effective tool at de-noising images. Indeed, our data suggests that applying a PCA pre-processor to a set of noised images can significantly improve data quality and aid in classification tasks, particularly when noise levels are significant.

Moreover, we find that while the output qualities of the three algorithms we analyzed, **NIPALS**, **Full SVD**, and **Simultaneous Iteration**, are quite similar, we find that their computational complexity differs greatly. For image de-noising, due to the large number of components that need to be calculated, **Full SVD** appears to be by far the lowest in cost. Nevertheless, our results suggest that this will not always be the case, and in problems with different data requirements, one of the other variants may be the most efficient.

Finally, we raise several possible avenues of further research. One of the main characteristics of our approach to PCA is that it makes no assumptions about the input data. In particular, our PCA filter is not specific to a single classifier and relies only on the noised data. Tailoring PCA to the specific classification task, perhaps by including the unnoised training data may lead to larger improvements in classifier performance. Furthermore, the improvements in performance we found when converting *NumPy* code to *Numba*, and the speedups obtained by utilizing a GPU suggest that there are many options to further optimize implementations of the studied algorithms.